
工研院產業學院
「AOI自動光學檢測瑕疵分類實作班」
(A)
亞洲大學朱學亭副教授
2020/07/01

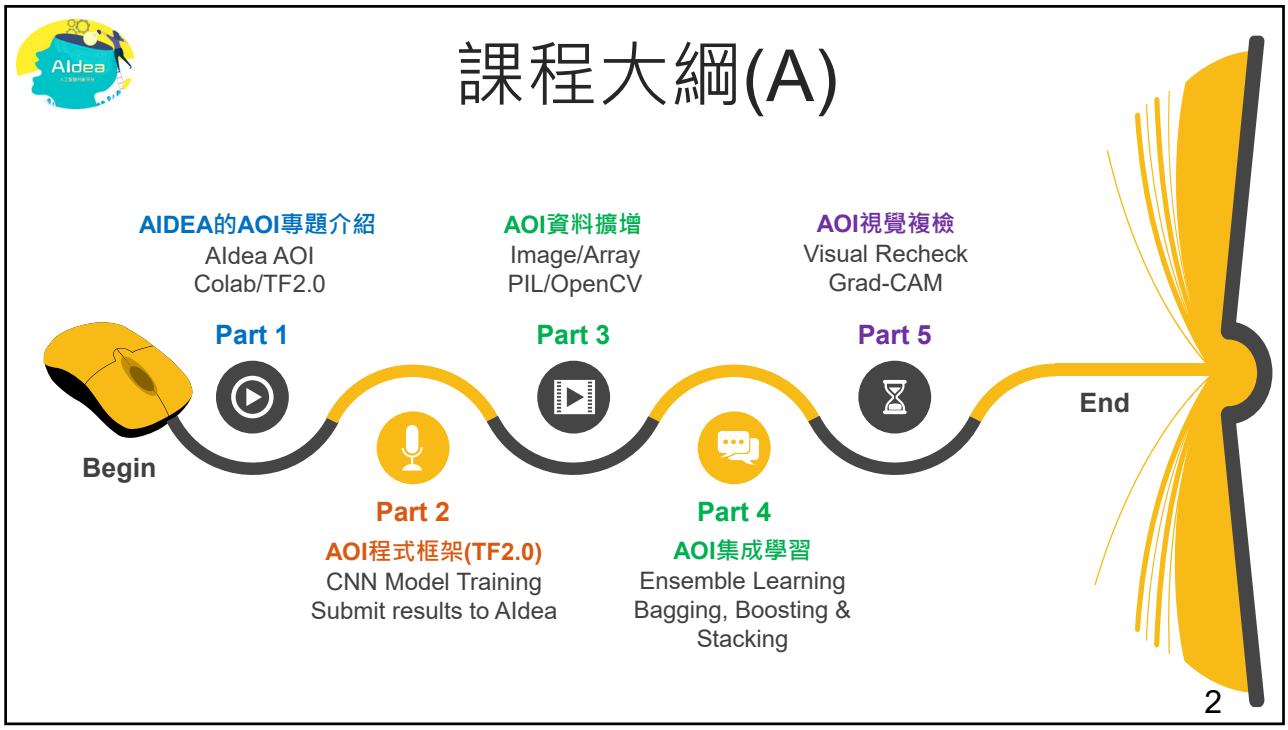
課程大綱(A)
AIDEA的AOI專題介紹
AIdea AOI
Colab/TF2.0
Part 1
AOI資料擴增
Image/Array
PIL/OpenCV
Part 3
AOI視覺複檢
Visual Recheck
Grad-CAM
Part 5
Begin
End
Part 2
AOI程式框架(TF2.0)
CNN Model Training
Submit results to AIdea
Part 4
AOI集成學習
Ensemble Learning
Bagging, Boosting &
Stacking
2

啟思博Kissipo(KISS + IPO)學深度學習

- **KISS principle**
  - "keep it simple, stupid" or "keep it stupid simple", is a design principle noted by the U.S. Navy in 1960.
  - https://en.wikipedia.org/wiki/KISS_principle
- **IPO model**
  - The input–process–output (IPO) model is a widely used approach in systems analysis and software engineering for describing the structure of an information processing program or other process.
  - https://en.wikipedia.org/wiki/IPO_model

3

線上工具:
Github, Colab, Google Forms

Github:

➢ https://github.com/htchu/aidea-auaoi

Colab:

✓ Google Colaboratory (Coding TensorFlow)
✓ Jupyter notebook environment

Google Forms

☐ 課前問卷： https://forms.gle/N3abrJR4BJXsyQXy9
☐ 課後問卷： https://forms.gle/JDCBbQhU4vQXqdfP8

Google Drive

model https://drive.google.com/file/d/____35v1fkLU3EC-xZCIq6mmvYkKyRUq9/view?usp=sharing
dataset https://drive.google.com/file/d/____y_6pkMwLrg05A4f8S5dRzyu4I5j8Q/view?usp=sharing

4

Github- AUAOI

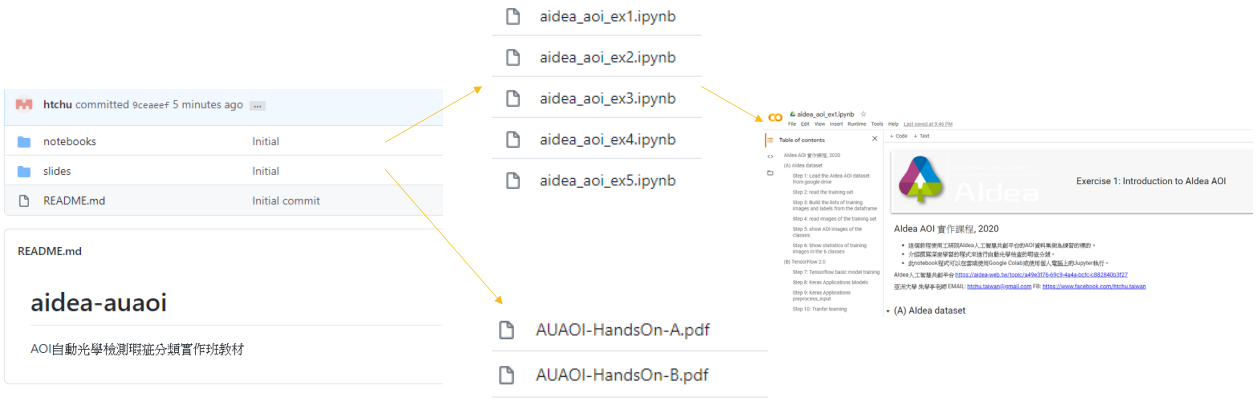

From Github to Google Colab

https://github.com/htchu/aidea-auaoi 5

Colab Limits

1. The GPUs available in Colab often include Nvidia K80s, T4s, P4s and P100s. There is no way to choose what type of GPU you can connect to in Colab at any given time.

2. Overall usage limits as well as idle timeout periods, maximum VM lifetime, GPU types available, and other factors vary over time.

6

1.AIdea的AOI專題介紹

7

AOI system architecture

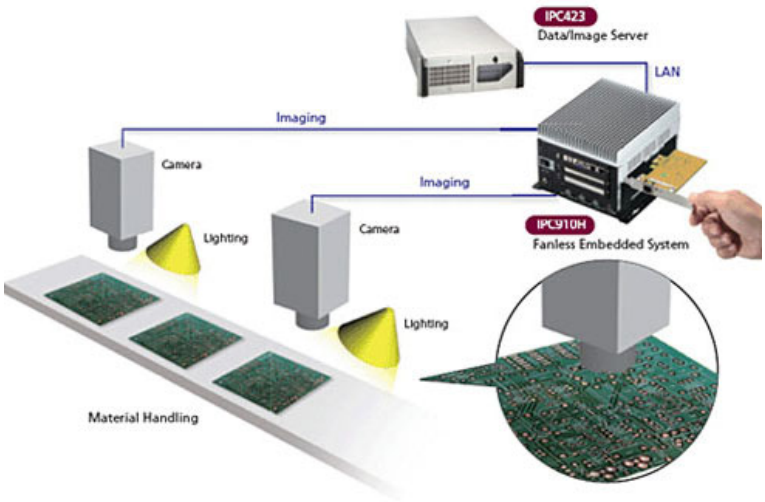

8

視覺檢測種類:SPI, AOI, AXI

SPI (Solder Paste Inspection)
AOI (Auto Optical Inspection)
AXI (Automatic X-ray Inspection)

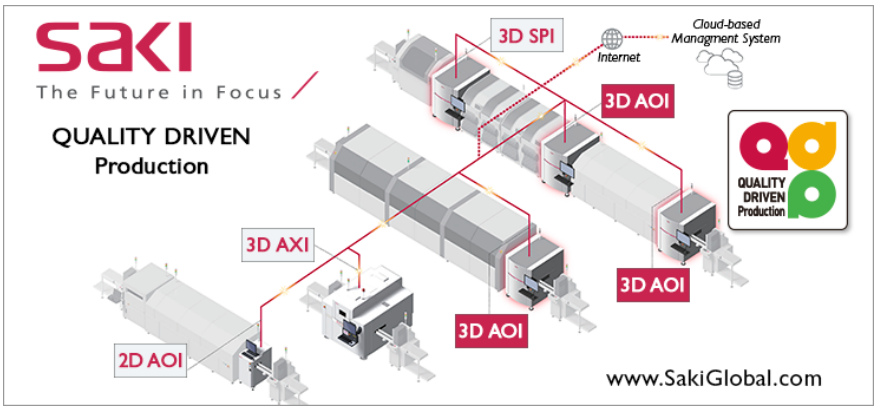

9

AOI工作流程差異

傳統機器學習的AOI[1990 年至今]

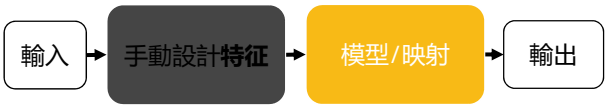

示例 [回歸和 SVM]

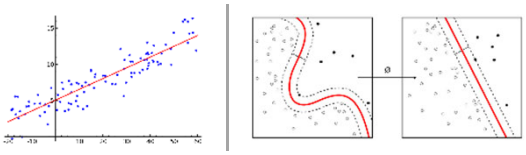

深度/端到端學習的AOI [2012 年至今]

示例 [卷積網路]

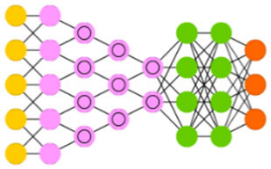

10

AOI with machine learning models

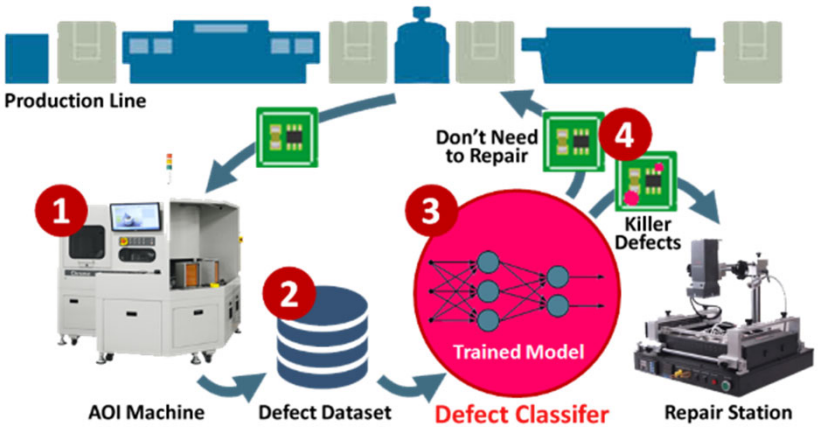

11

Classical ML-based AOI Machine

財團法人自強工業科學基金會
Tze-Chiang Foundation of Science & Technology

AOI自動光學檢測設備介紹與選用 如期開班

課程代碼：
02A040

上課時間：
2013/08/16(五)，9:00~16:00，共6小時

上課時數：
6 小時

課程大綱：
1.影像分析：
影像前處理：影像的運算、迴旋積運算、直方圖、直方圖修正、影像濾波處理
影像分割：二值化、多值化、自動閥值
影像後處理：邏輯運算、像素的近鄰、形態處理、骨架化、及物件標號

2. 影像量測：
邊界偵測：邊界、邊界點、梯度、一階導數運算子(Sobel、Prewitt、Robert)、二導數運算子(Laplacian、LOG)、鏈碼輪廓追蹤法
尺寸量測：最小平方法、霍式轉換法、直線搓合、圓及圓弧的搓合、線寬線距的量測、圓形物件的量測、矩形物體的量測。

3.影像辨識：
特徵抽取：特徵值(周長、離心率、真圓度、粗糙度、長短距離比、角數、平均距離、距離標準差、尤拉數、灰階標準差)
物件識別與分類:最近鄰分類器、K-NN分類器、類神經網路分類

12

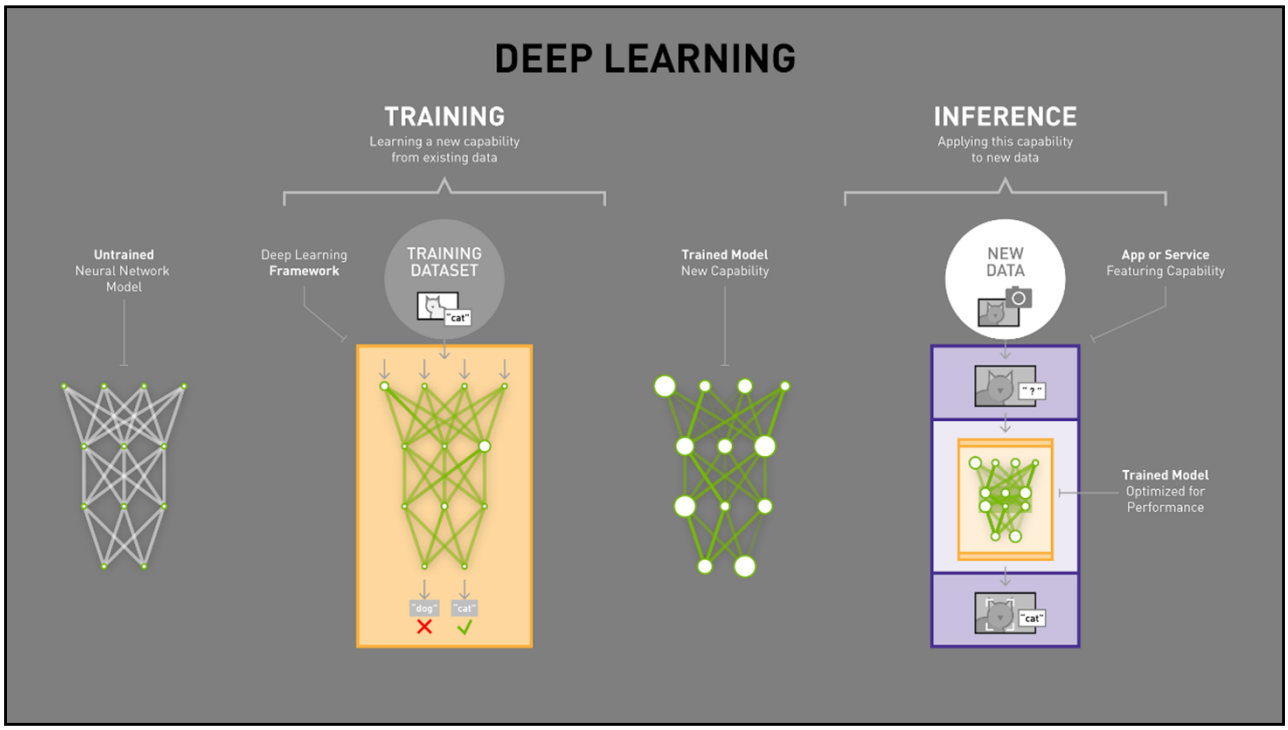

New DL-based AOI Machine

台達視覺檢測的 DAVS 以人工智慧為核心系統，結合既有的 AOI 系統，讓既有設備可延長使用年限以此保障製造業者過去的投資，而人工智慧與 AOI 整合的模式，也提升了產品的檢出率。

https://buzzorange.com/techorange/2019/09/05/delta-aoi-system/

14

深度學習方法 - 訓練

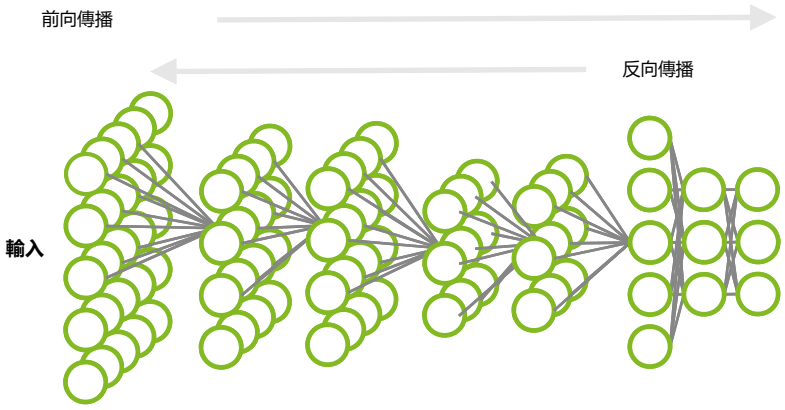

流程

- **前向傳播為每個訓練圖像生成一個推斷標籤**
- **使用損失函數計算每個圖像的已知標籤與預測標籤之間的差異**
- **在後向傳播過程中調整權重**
- **重複該流程**

15

Reducing Loss: An Iterative Approach

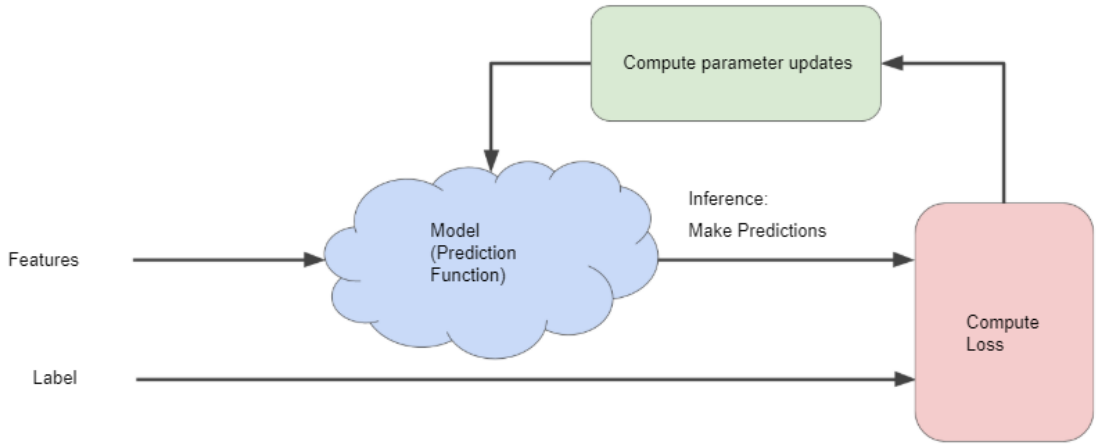

16

Convolutional neural network(CNN)

Add convolution and pooling layers before feedforward neural network

卷積計算
- 步伐(stride)
- 填充(padding)

池化運算

```
keras.layers.Conv2D(filters, kernel_size, strides=(1, 1),
padding='valid', data_format=None, dilation_rate=(1, 1),
activation=None, use_bias=True,
kernel_initializer='glorot_uniform', bias_initializer='zeros',
kernel_regularizer=None, bias_regularizer=None,
activity_regularizer=None, kernel_constraint=None,
bias_constraint=None)
```

```
keras.layers.MaxPooling2D(pool_size=(2, 2),
strides=None, padding='valid', data_format=None)
```

CNN History

楊立昆

Yann LeCun, Professor of Computer Science
The Courant Institute of Mathematical Sciences
New York University
Room 1220, 715 Broadway, New York, NY 10003, USA.
(212)998-3283 yann@cs.nyu.edu

In 1995, Yann LeCun and Yoshua Bengio introduced the concept of convolutional neural networks.

18

Convolutional layer

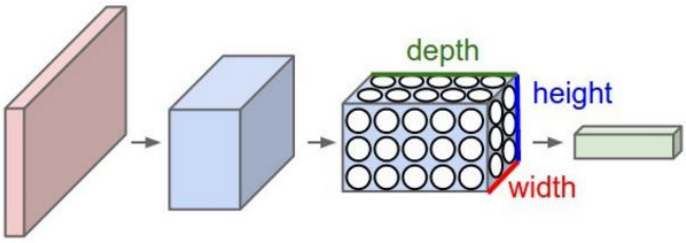

Pooling layer

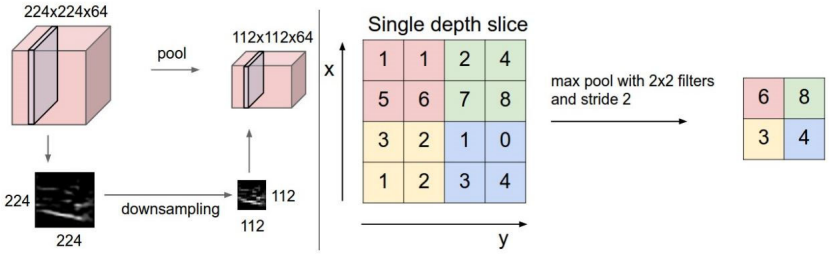

19

Keras

https://keras.io/zh/

François Chollet 弗朗索瓦·喬萊特
Deep learning @google
Creator of Keras, neural networks library.
Author of 'Deep Learning with Python'.

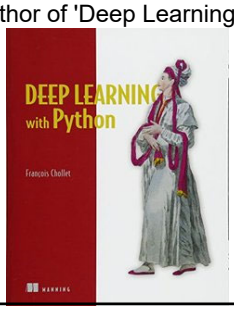

20

TensorFlow v2.2.0

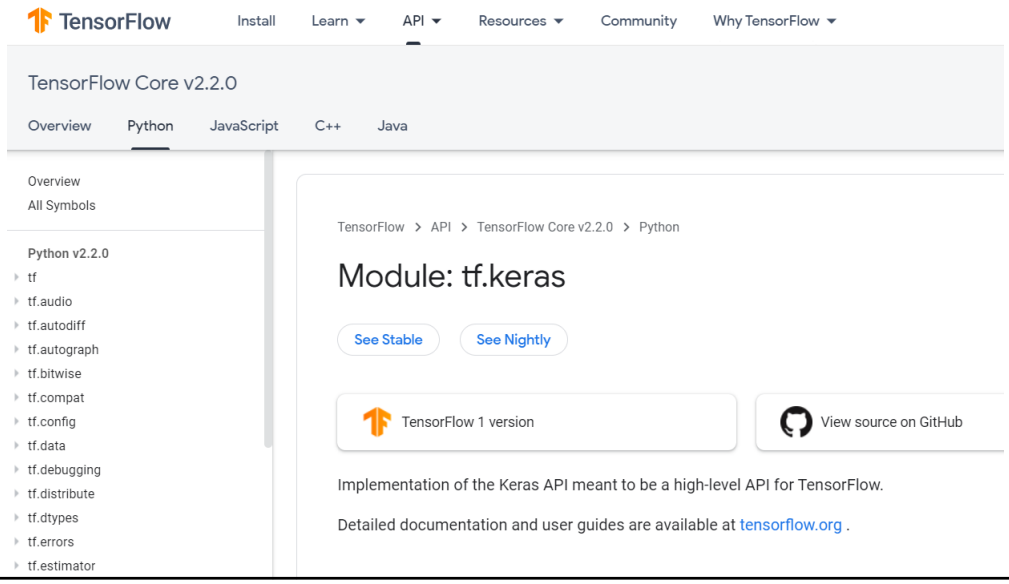

21

AIdea AOI Project

https://aidea-web.tw/topic/a49e3f76-69c9-4a4a-bcfc-c882840b3f27

22

Step 1: Load the dataset from google drive

AUAOI Ex1.

▾ Step 1: Load the AIdea AOI dataset from google drive

```
[ ] from google_drive_downloader import GoogleDriveDownloader
    GoogleDriveDownloader.download_file_from_google_drive(file_id='      yhVsQZU2iiK19xlJubw0afQ
```

```
Downloading      yhVsQZU2iiK19xlJubw0afQ2EMu5 into ./content... Done.
Unzipping...Done.
```

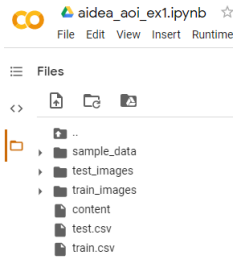

23

Step 2:

AUAOI Ex1.

Step 2: read the training set

train.csv

```
[2] import pandas as pd
    df_train = pd.read_csv("          ")
    print(df_train.shape)
```

```
(2528, 2)
```

24

Step 3: Build the lists of training images and labels


AUAOI Ex1.

```
train_files = df_train.iloc[:,].values
train_labels = df_train.iloc[:,].values
print(train_labels[:10])
```

•train_images.zip：訓練所需的影像資料（PNG格式），共計 2,528 張。
•train.csv：包含 2 個欄位，ID 和 Label。
- ID：影像的檔名。
- Label：瑕疵分類類別（0 表示 normal，1 表示 void，2 表示 horizontal defect，3 表示 vertical defect，4 表示 edge defect，5 表示 particle）。

•test_images.zip：測試所需的影像資料（PNG格式），共計 10,142 張。
•test.csv：包含 2 個欄位，ID 和 Label。
- ID：影像的檔名。
- Label：瑕疵分類類別（其值只能是下列其中之一：0、1、2、3、4、5）。

25

Step 4: read images of the training set


AUAOI Ex1.

train_images/

```
train_path ="      images/"
train_images = []
from tensorflow.keras.preprocessing import image
for file in train_files:
    img = image.load_img(train_path+file, color_mode="rgb", target_size = (299, 299))
    train_images.
    if len(train_images)%100 == 0:
      print('.', end='')
print(len(train_images))
```

append(img)

```
.........................2528
```

numpy.ndarray（值0-255）

```
from tensorflow.keras.preprocessing import image
img = image.load_img(file, color_mode="rgb", target_size = (299, 299))
```

```
from PIL import Image
img = Image.open(file)

import matplotlib.pyplot as plt
img = plt.imread(file)
```

```
from skimage import io
img = io.imread(file)
import cv2
img = cv2.imread(file) #OpenCV is BGR, Pillow is RGB
#imread is deprecated in SciPy 1.0.0
```

26

Step 5: show AOI images of the classes

AUAOI Ex1.

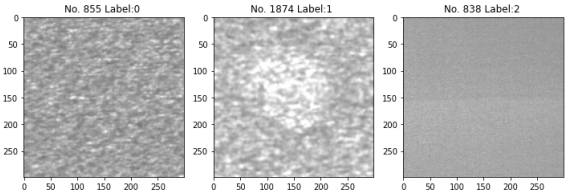

0 (normal), 1 (void), 2 (horizontal defect)

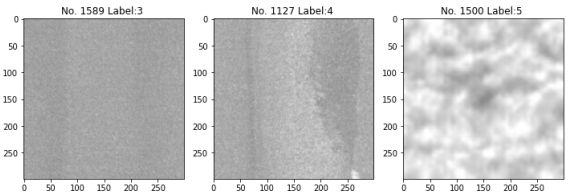

3 (vertical defect), 4 (edge defect), 5 (particle)

Step 6: Show statistics of training images in the 6 classes

AUAOI Ex1.

train_labels

```
import numpy as np
labels, counts = np.unique(                , return_counts=True)
print(labels, counts)
```

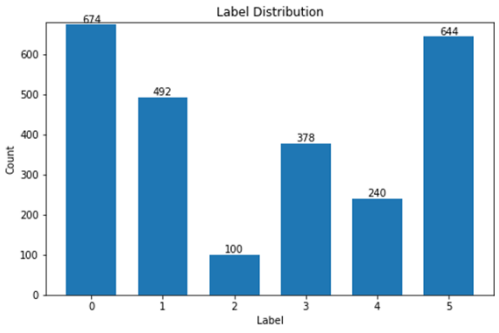

Step 7a: Tensorflow basic model training

AUAOI Ex1.

```
[6] import tensorflow as tf
    print(tf.__version__)
    print(tf.config.list_physical_devices('GPU'))
```

```
2.2.0
[PhysicalDevice(name='/physical_device:GPU:0', device_type='GPU')]
```

```
[7] !nvidia-smi
```

```
Sun Jun 28 17:39:16 2020
+-----------------------------------------------------------------------------+
| NVIDIA-SMI 450.36.06    Driver Version: 418.67       CUDA Version: 10.1     |
|-------------------------------+----------------------+----------------------+
| GPU  Name        Persistence-M| Bus-Id        Disp.A | Volatile Uncorr. ECC |
| Fan  Temp  Perf  Pwr:Usage/Cap|         Memory-Usage | GPU-Util  Compute M. |
|                               |                      |               MIG M. |
|===============================+======================+======================|
|   0  Tesla P100-PCIE...  Off  | 00000000:00:04.0 Off |                    0 |
| N/A   39C    P0    25W / 250W |     10MiB / 16280MiB |      0%      Default |
|                               |                      |                 ERR! |
+-------------------------------+----------------------+----------------------+
```


Step 7b:

AUAOI Ex1.

```
[8] mnist = tf.keras.datasets.mnist
    (x_train, y_train), (x_test, y_test) = mnist.load_data()
    x_train, x_test = [redacted] / 255.0, [redacted] / 255.0
```

```
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz
11493376/11490434 [==============================] - 0s 0us/step
```

```
[9] x_train = x_train[..., tf.[redacted]]
    x_test  = x_test[..., tf.[redacted]]
    print(x_train.shape)
    print(y_train.shape)
```

```
(60000, 28, 28, 1)
(60000,)
```

tf.newaxis

Step 7c:


```
[10] model = tf.keras.models.Sequential([
        tf.keras.layers.Conv2D(filters = 32, kernel_size = (3,3),padding = 'Same',
                    activation ='relu', input_shape = (28,28,1)),
        tf.keras.layers.Conv2D(filters = 64, kernel_size = (3,3),padding = 'Same',
                    activation ='relu'),
        tf.keras.layers.MaxPool2D(pool_size=(2,2)),
        tf.keras.layers.Dropout(0.25),
        tf.keras.layers.Flatten(),
        tf.keras.layers.Dense(128, activation='relu'),
        tf.keras.layers.Dropout(0.5),
        tf.keras.layers.Dense(10, activation='       ')
    ])
```

softmax

```
[11] opt = tf.keras.optimizers.SGD(learning_rate=0.01) #lr =0.01
     loss_fn = tf.keras.losses.      CategoricalCrossentropy(from_logits=True)
```

SparseCategoricalCrossentropy

```
[ ] model.compile(optimizer=opt, loss=loss_fn, metrics=['accuracy'])
```


Step 7d:


```
[13] model.     (x_train, y_train, epochs=5)

Epoch 1/5
1875/1875 [==============================] - 6s 3ms/step - loss: 1.9838 - accuracy: 0.5317
Epoch 2/5
1875/1875 [==============================] - 6s 3ms/step - loss: 1.6146 - accuracy: 0.8565
Epoch 3/5
1875/1875 [==============================] - 6s 3ms/step - loss: 1.5843 - accuracy: 0.8836
Epoch 4/5
1875/1875 [==============================] - 6s 3ms/step - loss: 1.5690 - accuracy: 0.8963
Epoch 5/5
1875/1875 [==============================] - 6s 3ms/step - loss: 1.5581 - accuracy: 0.9075
<tensorflow.python.keras.callbacks.History at 0x7f696002c0b8>
```

fit

```
[14] model.       (x_test,  y_test, verbose=1)

313/313 [==============================] - 1s 2ms/step - loss: 1.5203 - accuracy: 0.9435
[1.5202512741088867, 0.9434999823570251]
```

evaluate

```
y_predicts = model.       (x_test)
y_predicts[0]

array([1.4306398e-19, 2.1704624e-23, 3.9827185e-17, 1.3840105e-15,
       1.0869419e-19, 1.6800059e-18, 3.7071232e-24, 1.0000000e+00,
       6.0183126e-16, 4.9968688e-13], dtype=float32)
```

predict

Step 8: Keras Applications Models


```
from tensorflow.keras.applications import InceptionV3
from tensorflow.keras.applications.inception_v3 import preprocess_input
model = InceptionV3(include_top = True, input_shape=(299,299,3), weights=None, classes=num_
```

```
from tensorflow.keras.applications import Xception
from tensorflow.keras.applications.xception import preprocess_input
model = Xception(include_top = True, input_shape=(299,299,3), weights=None, classes=num_cla
```

```
from tensorflow.keras.applications import NASNetLarge
from tensorflow.keras.applications.nasnet import preprocess_input
model = NASNetLarge(include_top = True, input_shape=(299,299,3), weights=None, classes=num_
```

```
from tensorflow.keras.applications import InceptionResNetV2
from tensorflow.keras.applications.inception_resnet_v2 import preprocess_input
model = InceptionResNetV2(include_top = True, input_shape=(299,299,3), weights=None, classe
```

```
from tensorflow.keras.applications import MobileNetV2
from tensorflow.keras.applications.mobilenet_v2 import preprocess_input
model = MobileNetV2(include_top = True, input_shape=(299,299,3), weights=None, classes=num_
```

```
from tensorflow.keras.applications import ResNet50V2
from tensorflow.keras.applications.resnet_v2 import preprocess_input
model = ResNet50V2(include_top = True, input_shape=(299,299,3), weights=None, classes=num_c
```

Model	Size	Top-1 Accuracy	Top-5 Accuracy	Parameters	Depth
Xception	88 MB	0.790	0.945	22,910,480	126
VGG16	528 MB	0.713	0.901	138,357,544	23
VGG19	549 MB	0.713	0.900	143,667,240	26
ResNet50	98 MB	0.749	0.921	25,636,712	-
ResNet101	171 MB	0.764	0.928	44,707,176	-
ResNet152	232 MB	0.766	0.931	60,419,944	-
ResNet50V2	98 MB	0.760	0.930	25,613,800	-
ResNet101V2	171 MB	0.772	0.938	44,675,560	-
ResNet152V2	232 MB	0.780	0.942	60,380,648	-
InceptionV3	92 MB	0.779	0.937	23,851,784	159
InceptionResNetV2	215 MB	0.803	0.953	55,873,736	572
MobileNet	16 MB	0.704	0.895	4,253,864	88
MobileNetV2	14 MB	0.713	0.901	3,538,984	88
DenseNet121	33 MB	0.750	0.923	8,062,504	121
DenseNet169	57 MB	0.762	0.932	14,307,880	169
DenseNet201	80 MB	0.773	0.936	20,242,984	201
NASNetMobile	23 MB	0.744	0.919	5,326,716	-
NASNetLarge	343 MB	0.825	0.960	88,949,818	-
EfficientNetB0	29 MB	-	-	5,330,571	-
EfficientNetB1	31 MB	-	-	7,856,239	-
EfficientNetB2	36 MB	-	-	9,177,569	-
EfficientNetB3	48 MB	-	-	12,320,535	-
EfficientNetB4	75 MB	-	-	19,466,823	-
EfficientNetB5	118 MB	-	-	30,562,527	-
EfficientNetB6	166 MB	-	-	43,265,143	-
EfficientNetB7	256 MB	-	-	66,658,687	-

33

Step 9: Keras Applications preprocess_input


```
[16] from tensorflow.keras.preprocessing.image import img_to_array
     from tensorflow.python.keras.applications.imagenet_utils import preprocess_input
     x = image.img_to_array(train_images[0])
     img_array = preprocess_input(x, mode = 'tf' )
     print(img_array[0 , 0 , 0])

     0.3411765
```

```
[17] from tensorflow.keras.preprocessing.image import img_to_array
     from tensorflow.python.keras.applications.imagenet_utils import preprocess_input
     x = image.img_to_array(train_images[0])
     img_array = preprocess_input(x, mode = 'torch' )
     print(img_array[0 , 0 , 0])

     0.810429
```

```
     from tensorflow.keras.preprocessing.image import img_to_array
     from tensorflow.python.keras.applications.imagenet_utils import preprocess_input
     x = image.img_to_array(train_images[0])
     img_array = preprocess_input(x, mode = 'caffe' )
     print(img_array[0 , 0 , 0])

     67.061
```

	Input size	Data format	mode
Xception	299x299	channels_last	tf
VGG16	224x224	channels_first / channels_last	caffe
VGG19	224x224	channels_first / channels_last	caffe
ResNet50	224x224	channels_first / channels_last	caffe
InceptionV3	299x299	channels_first / channels_last	tf
InceptionResNetV2	299x299	channels_first / channels_last	tf
MobileNet	224x224	channels_last	tf
DenseNet	224x224	channels_first / channels_last	torch
NASNet	331x331 / 224x224	channels_first / channels_last	tf
MobileNetV2	224x224	channels_last	tf

*mode = **caffe***
(will convert the images from RGB to BGR, then will zero-center each color channel with respect to the ImageNet dataset)
減去ImageNet平均 BGR [103.939, 116.779, 123.68]
*mode = **tf***
(will scale pixels between -1 and 1) **除以127.5，然後減 1。**
*mode = **torch***
(will scale pixels between 0 and 1 and then will normalize each channel with respect to the ImageNet dataset)
除以255，減去ImageNet平均[0.485, 0.456, 0.406] ，除以標準差[0.229, 0.224, 0.225]。

34

Step 10: Tranfer learning

AUAOI Ex1.

```
#the InceptionV3 model
num_classes = 6
from tensorflow.keras.applications import InceptionV3
base_model = InceptionV3(include_top = False, input_shape=(299,299,3), weights='imagenet', classes=num_classes)
base_model.summary()
```

```
base_model.trainable = False
last_layer = base_model.output
last_layer=Flatten()(last_layer)
last_layer=Dropout(0.3)(last_layer)
out = Dense(num_classes, activation='softmax', name='softmax')(last_layer)
custom_model = Model(base_model.input, out)
model.summary()
```

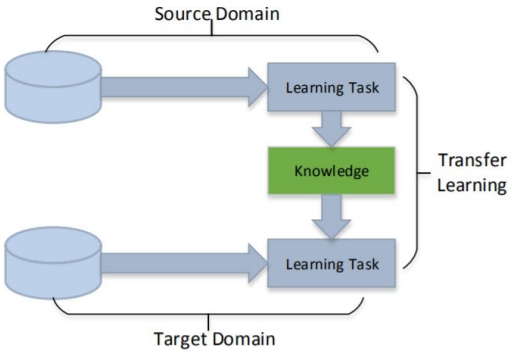

https://www.coderbridge.com/@gueiyajhang/54584ea6d4c240aeb3b8ae4af3a0531a

35

2. TF 2.0 AOI程式框架

aidea_aoi_ex2.ipynb

File Edit View Insert Runtime Tools Help All changes saved

Table of contents

- Exercise 2: Simple solution
- **Step 1: Load the dataset from google drive**
- Step 2: Import python libraries
- Step 3: read the training set
- Step 4: Show statistics of training images
- Step 5: Choose one of CNN models
- Step 6: Instancing an ImageDataGenerator
- Step 7: Set up a train_generator with flow_from_dataframe
- Step 8: step_size_train
- Step 9: ModelCheckpoint
- Step 10: EarlyStopping
- Step 11: Compile model
- Step 12: Train model
- Step 13: Evaluate saved checkpoints
- Step 14: Save the trained model
- Step 15: Check training results
- Step 16: Analyze training results
- Step 17: Load the test set
- Step 18: Set up a test_generator with flow_from_dataframe
- Step 19: step_size_train
- Step 20: Check test results
- Step 21: Output test results

+ Code + Text

ARTIFICIAL INTELLIGENCE COLLABORATION PLATFORM Aidea

Exercise 2: Basic AOI algorithm

Exercise 2: Simple solution

- single CNN model
- ImageDataGenerator
- ModelCheckpoint
- EarlyStopping
- Submit results

Step 1: Load the dataset from google drive

36

Step 1: Load the dataset from google drive

AUAOI Ex2.

▾ Step 1: Load the AIdea AOI dataset from google drive

```
[ ] from google_drive_downloader import GoogleDriveDownloader
    GoogleDriveDownloader.download_file_from_google_drive(file_id='      yhVsQZU2iiK19xlJubw0afQ
```

```
Downloading      yhVsQZU2iiK19xlJubw0afQ2EMu5 into ./content... Done.
Unzipping...Done.
```

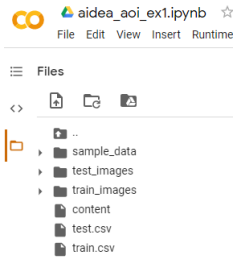

37

Step 2: Import python libraries

AUAOI Ex2.

```
[2] import tensorflow as tf
    tf.config.experimental.set_memory_growth(tf.config.list_physical_devices('GPU')[0], True)
    print(tf.__version__)
    print(tf.config.list_physical_devices('GPU'))
```

```
2.2.0
[PhysicalDevice(name='/physical_device:GPU:0', device_type='GPU')]
```

```
[3] import numpy as np
    import matplotlib.pyplot as plt
    %matplotlib inline
```

```
[4] from tensorflow.keras import Sequential
    from tensorflow.keras.models import Model
    from tensorflow.keras.layers import Dense, Activation, Flatten
    from tensorflow.keras.layers import Input
    from tensorflow.keras.layers import Dropout, Flatten, Activation
    from tensorflow.keras.layers import Conv2D, MaxPooling2D
    from tensorflow.keras.optimizers import Adam,SGD,Adagrad,Adadelta,RMSprop
```

38

Step 3: read the training set

AUAOI Ex2.

train.csv str

```
import pandas as pd
df_train = pd.read_csv("[redacted]",dtype=[redacted])
print(df_train.shape)
```

```
[6] df_train.head()
```

	ID	Label
0	train_00000.png	0
1	train_00001.png	1
2	train_00002.png	1
3	train_00003.png	5
4	train_00004.png	5

```
[7] train_files = df_train.iloc[:,0].values
    train_labels = df_train.iloc[:,1].values
    print(train_labels[:10])
```

```
['0' '1' '1' '5' '5' '5' '3' '0' '3' '5']
```

39

Step 4: Show statistics of training images

AUAOI Ex2.

```
import seaborn as sns
g = sns.countplot(train_labels)
```

```
/usr/local/lib/python3.6/dist-packages/statsmodels/tools/_te
  import pandas.util.testing as tm
```

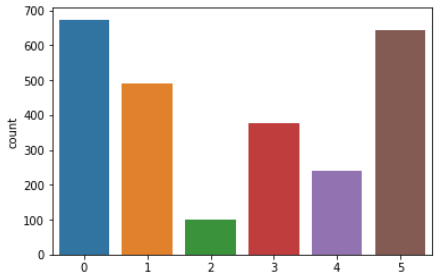


```
num_classes=6
```

40

Step 5: Choose one of CNN models

AUAOI Ex2.

```
[11] from tensorflow.keras.applications import
     model = (include_top = True, input_shape=(299,299,3), weights=None, classes=num_classes)

[12] model.summary()
```

- DenseNet121(...): Instantiates the Densenet121 architecture.
- DenseNet169(...): Instantiates the Densenet169 architecture.
- DenseNet201(...): Instantiates the Densenet201 architecture.
- InceptionResNetV2(...): Instantiates the Inception-ResNet v2 architecture.
- InceptionV3(...): Instantiates the Inception v3 architecture.
- MobileNet(...): Instantiates the MobileNet architecture.
- MobileNetV2(...): Instantiates the MobileNetV2 architecture.
- NASNetLarge(...): Instantiates a NASNet model in ImageNet mode.
- NASNetMobile(...): Instantiates a Mobile NASNet model in ImageNet mode.
- ResNet101(...): Instantiates the ResNet101 architecture.
- ResNet101V2(...): Instantiates the ResNet101V2 architecture.
- ResNet152(...): Instantiates the ResNet152 architecture.
- ResNet152V2(...): Instantiates the ResNet152V2 architecture.
- ResNet50(...): Instantiates the ResNet50 architecture.
- ResNet50V2(...): Instantiates the ResNet50V2 architecture.
- VGG16(...): Instantiates the VGG16 model.
- VGG19(...): Instantiates the VGG19 architecture.
- Xception(...): Instantiates the Xception architecture.

41

Step 6: Instancing an ImageDataGenerator

AUAOI Ex2.

preprocess_input

```
[13] from tensorflow.keras.preprocessing.image import ImageDataGenerator
     from tensorflow.keras.applications.xception import preprocess_input
     img_gen = ImageDataGenerator(preprocessing_function=
```

```
tf.keras.preprocessing.image.ImageDataGenerator(
    featurewise_center=False, samplewise_center=False,
    featurewise_std_normalization=False, samplewise_std_normalization=False,
    zca_whitening=False, zca_epsilon=1e-06, rotation_range=0, width_shift_range=0.0,
    height_shift_range=0.0, brightness_range=None, shear_range=0.0, zoom_range=0.0,
    channel_shift_range=0.0, fill_mode='nearest', cval=0.0, horizontal_flip=False,
    vertical_flip=False, rescale=None, preprocessing_function=None,
    data_format=None, validation_split=0.0, dtype=None
)
```

https://www.tensorflow.org/api_docs/python/tf/keras/preprocessing/image/ImageDataGenerator

42

Step 7: Set up a train_generator with flow_from_dataframe

AUAOI Ex2.

```
[14] train_generator = img_gen.flow_from_dataframe(dataframe=[redacted],
                directory="[redacted]",
                x_col="[redacted]",
                y_col="[redacted]",
                subset=None,
                batch_size=8,
                shuffle=False,
                class_mode="categorical",
                color_mode="rgb",
                target_size=(299,299))
```

```
Found 2528 validated image filenames belonging to 6 classes.
```

```
[15] train_generator.class_indices
```

```
{'0': 0, '1': 1, '2': 2, '3': 3, '4': 4, '5': 5}
```

`df_train`

`train_images`

`ID`

`Label`

```
flow_from_dataframe(
   dataframe, directory=None, x_col='filename', y_col='class', weight_col=None,
   target_size=(256, 256), color_mode='rgb', classes=None,
   class_mode='categorical', batch_size=32, shuffle=True, seed=None,
   save_to_dir=None, save_prefix='', save_format='png', subset=None,
   interpolation='nearest', validate_filenames=True, **kwargs
)
```

https://www.tensorflow.org/api_docs/python/tf/keras/preprocessing/image/ImageDataGenerator

43

Step 8: step_size_train

AUAOI Ex2.

```
if train_generator.n % train_generator.batch_size ==0:
  step_size_train=train_generator.n//train_generator.batch_size
else:
  step_size_train=train_generator.n//train_generator.batch_size + 1
print(step_size_train)
```

44

Step 9:

ModelCheckpoint

Callback to save the Keras model or model weights at some frequency.

AUAOI Ex2.

```
[17] # Include the epoch in the file name (uses `str.format`)
     import os
     checkpoint_path = "training_cp/cp-{epoch:03d}.ckpt"
     checkpoint_dir = os.path.dirname(checkpoint_path)
     # Create a callback that saves the model's weights
     cp_callback = tf.keras.callbacks.ModelCheckpoint(filepath=[          ], save_weights_only=True)
```

checkpoint_path

```
tf.keras.callbacks.ModelCheckpoint(
    filepath, monitor='val_loss', verbose=0, save_best_only=False,
    save_weights_only=False, mode='auto', save_freq='epoch', **kwargs
)
```

https://www.tensorflow.org/api_docs/python/tf/keras/callbacks/ModelCheckpoint

45

Step 10:

EarlyStopping

Stop training when a monitored metric has stopped improving.

AUAOI Ex2.

```
[18] # Create a callback that stop the model
     es_callback = tf.keras.callbacks.EarlyStopping(monitor='[    ]', patience=5)
```

loss

```
tf.keras.callbacks.EarlyStopping(
    monitor='val_loss', min_delta=0, patience=0, verbose=0, mode='auto',
    baseline=None, restore_best_weights=False
)
```

46

Step 11: Compile model


```
[19] #compile model using accuracy to measure model performance
     from tensorflow.keras import optimizers
     model.compile(loss='categorical_crossentropy',
                   optimizer=optimizers.[    ](lr=3e-3),
                   metrics=['accuracy'])
```

Adam/SGD

```
compile(
    optimizer='rmsprop', loss=None, metrics=None, loss_weights=None,
    sample_weight_mode=None, weighted_metrics=None, **kwargs
)
```

tf.keras.optimizers

Optimizer	特點
SGD	• 有機會跳出目前局部收斂進而進到另一個局部收斂而得到最小值，而得到全局最小值 • 需自行設定learning rate，較難選擇到合適的learning rate • 會造成loss function有嚴重的震蕩 • 需要較長時間收斂至最小值
Momentum	• 能夠在相關方向加速SGD，抑制SGD的嚴重震蕩，進而加快收斂 • 需自行設定learning rate與β，有可能會使參數的移動方向偏移梯度下分的方向，進而導至沒有那麼快速的收斂
AdaGrad	• 能夠自動調整learning rate，進而調整收斂 • 適合處理稀疏梯度 • 依然需要人工設置一個全局的learning rate • 後期，分母梯度平方的累加會越來越大，會使梯度趨近於0，使得訓練結束
Adam	• 結合了AdaGrad與Momentum的優點 • 適用於大數據集和高維空間的資料 • 目前最常使用的一個Optimizer

47

Step 12: Train model

cp_callback, es_callback

train_generator step_size_train

```
hist = model.fit_generator(generator=[    ], steps_per_epoch=[    ], callbacks=[[    ], [    ]], epochs=100)
```

```
WARNING:tensorflow:From <ipython-input-20-5be68c6e0f2d>:1: Model.fit_generator (from tensorflow.python.keras.engine.training) is deprecate
Instructions for updating:
Please use Model.fit, which supports generators.
Epoch 1/100
111/316 [=========>.....................] - ETA: 1:09 - loss: 1.7552 - accuracy: 0.4155
```

```
fit_generator(
    generator, steps_per_epoch=None, epochs=1, verbose=1, callbacks=None,
    validation_data=None, validation_steps=None, validation_freq=1,
    class_weight=None, max_queue_size=10, workers=1, use_multiprocessing=False,
    shuffle=True, initial_epoch=0
)
```

48

Step 13: Evaluate saved checkpoints

AUAOI Ex2.

```
##checkpoint 1
model.load_weights("training_cp/cp-001.ckpt")
train_generator.reset()
model.evaluate_generator(generator=train_generator, steps=step_size_train, verbose=1)
```

```
<tensorflow.python.training.tracking.util.CheckpointLoadStatus at 0x7fcc08582cc0>
```

```
##checkpoint 2
model.load_weights("training_cp/cp-001.ckpt")
train_generator.reset()
model.evaluate_generator(generator=train_generator, steps=step_size_train, verbose=1)
```

```
##checkpoint 3
model.load_weights("training_cp/cp-001.ckpt")
train_generator.reset()
model.evaluate_generator(generator=train_generator, steps=step_size_train, verbose=1)
```

49

Step 14: Save the trained model

AUAOI Ex2.

```
model.load_weights("training_cp/cp-001.ckpt")
model.save("AOI-InceptionV3-0626.h5")
```

```
save(
    filepath, overwrite=True, include_optimizer=True, save_format=None,
    signatures=None, options=None
)
```

https://www.tensorflow.org/api_docs/python/tf/keras/Model#save

50

深度學習模型輸出

- model.save(file)
- tf.saved_model.save(model, path)

- ckpt = tf.train.Checkpoint()
- ckpt.save()

51

TF深度學習模型輸入

- model=tf.keras.models.load_model (file)
- model=tf.saved_model.load(path)
- ckpt = tf.train.Checkpoint()
- ckpt.restore()

TF深度學習模型輸出

- model.save(file)
- tf.saved_model.save(model, path)
- ckpt = tf.train.Checkpoint()
- ckpt.save()

52

Step 15: Check training results

AUAOI Ex2.

train_generator

```
#y_predictions = model.predict(X_train, batch_size=20)
train_generator.reset()
y_predictions = model.predict_generator(generator=[redacted], steps=step_size_train, verbose=1)
```

```
WARNING:tensorflow:From <ipython-input-18-9c359a3ebada>:3: Model.predict_generator (from tensorflow.pytho
Instructions for updating:
Please use Model.predict, which supports generators.
316/316 [==============================] - 9s 29ms/step
```

```
print(y_predictions[:2])
type(y_predictions)
```

```
predicts = np.argmax(y_predictions,axis=1)
print(predicts[0:10])
```

```
[0 1 1 5 5 5 3 0 3 5]
```

```
labels = train_labels.astype(int)
print(labels[:10])
```

```
[0 1 1 5 5 5 3 0 3 5]
```

53

Step 16: Analyze training results

AUAOI Ex2.

labels predicts

```
from sklearn.metrics import confusion_matrix
confusion=confusion_matrix([redacted], [redacted])
print(confusion)
```

```
[[674   0   0   0   0   0]
 [  4 484   0   2   1   1]
 [  0   0 100   0   0   0]
 [  0   0   0 376   1   1]
 [  0   0   0   2 238   0]
 [  0   4   0   0   0 640]]
```

```
overkill= []
underkill = []
for i in range(train_num):
  if labels[i] == 0 and predicts[i] !=0:
    overkill.append(i)
  if labels[i] != 0 and predicts[i] ==0:
    underkill.append(i)
print('# of overkill= {}; # of underkill= {} '.format(len(overkill), len(underkill)))
```

54

Step 17: Load the test set

AUAOI Ex2.

```
df_test = pd.read_csv("[redacted]",dtype=str)
print(df_test.shape)
```

test.csv

```
df_test.head()
```

```
test_files  = df_test.iloc[:,0].values
test_labels = df_test.iloc[:,1].values
print(test_labels[:10])
```

55

Step 18: Set up a test_generator with flow_from_dataframe

AUAOI Ex2.

```
test_generator = img_gen.flow_from_dataframe(dataframe=[redacted]
            directory="[redacted]",
            x_col="ID",
            y_col="Label",
            batch_size=32,
            shuffle=False,
            class_mode=[redacted],
            target_size=(299,299))
```

df_test

test_images

None

56

Step 19: step_size_test

AUAOI Ex2.

```
if test_generator.n % test_generator.batch_size ==0:
  step_size_test=test_generator.n//test_generator.batch_size
else:
  step_size_test=test_generator.n//test_generator.batch_size + 1
print(step_size_test)
```

57

Step 20: Check test results

AUAOI Ex2.

```
#y_predictions = model.predict(X_train, batch_size=20)
test_generator.reset()
y_predictions = model.predict_generator(generator=test_generator, steps=step_size_test,verbose=1)
```

```
import numpy as np
predicts=np.argmax(y_predictions,axis=1)
predicts[:10]
```

58

Step 21:


AUAOI Ex2.

```
df_out = pd.DataFrame(df_test)
df_out.shape
```

```
df_out['Label'] = predicts
df_out.to_csv("0626-xception.csv", index=False)
```

59

~~End~~

Thanks!

Q&A

(1)Overfitting problem
(2)Training set/Validation set

行走江湖難免遇到八鴿...